2. DT KNN

Decision Tree

Impurity Measures:
- Misclassification rate: $i_E(t) = 1 - \max_c \pi_c$
- Entropy (Shannon): $i_H(t) = -\sum_{c_i} \pi_{ci} \log_2 \pi_{ci}$
- Gini index: $i_G(t) = 1 - \sum_{c_i} \pi_{ci}^2$

Greedy Optimization: Use $i_H(t)$ or $i_G(t)$ (not $i_E(t)$, which doesn't decrease impurity)

Information Gain:
- $p_L = \frac{N_{left}}{N}$ $\quad p_R = \frac{N_{right}}{N}$
- $\Delta i = i(t) - p_L \cdot i(t_L) - p_R \cdot i(t_R)$

Stopping Conditions: $i(t) = 0$ / $dmax$ / $N_{node} < tn$ / $\Delta i(s,t) < ts$

LOOCV: Equivalent to N-fold cross-validation

KNN

Prediction: $\hat{y} = \arg\max_c \sum_{x_i \in N_k(x)} \mathbb{I}(y_i = c)$

Distance Metrics:
- L1 (Manhattan): $d(x_1, x_2) = \sum_d |x_{1d} - x_{2d}|$
- L2 (Euclidean): $d(x_1, x_2) = \sqrt{\sum_d (x_{1d} - x_{2d})^2}$
- $L\infty$: $d(x_1, x_2) = \max_d |x_{1d} - x_{2d}|$
- Cosine Similarity: $\text{Sim}(x_1, x_2) = \frac{x_1^T x_2}{\|x_1\|\|x_2\|}$
- Mahalanobis Distance: $\sqrt{(x_1 - x_2)^T \Sigma^{-1} (x_1 - x_2)}$ (Σ positive semi-definite, symmetric)

Weighted KNN (inverse distance weighting, closer points more important):
- $\hat{y} = \arg\max_c \frac{1}{k} \sum_{x_i \in N_k(x)} \frac{1}{d(x_i, x)} \mathbb{I}(y_i = c)$

Hyperparameter Selection:
- k small → overfitting
- k large → underfitting
- Use odd number to avoid ties

Scale Issue: Normalization $x_f = \frac{x_i - \mu_i}{\sigma_i}$ (or use weighted distance)

Confusion Matrix

Ground \ Predict	1	0
1	TP	FN
0	FP	TN

Metrics:
- Precision: $\frac{TP}{TP+FP}$
- Sensitivity/Recall: $\frac{TP}{TP+FN}$
- Accuracy: $\frac{TP+TN}{TP+TN+FP+FN}$
- F1 Score: $\frac{2 \cdot \text{prec} \cdot \text{rec}}{\text{prec} + \text{rec}}$

3. Prob Method

Probabilistic Inference

Maximum Likelihood Estimation (MLE):
- $\theta_{MLE} = \arg\max_\theta p(D|\theta)$
- $p(D|\theta) = \prod_i^N p(x_i|\theta)$
- $E_{MLE} = -\ln p(D|\theta) = -\sum_i^N \ln p(x_i|\theta)$
- $\theta_{MLE} = \frac{|T|}{|T|+|H|}$

Maximum A Posteriori (MAP):
- $\theta_{MAP} = \arg\max_\theta p(\theta|D)$
- $p(\theta|D) \propto p(D|\theta)p(\theta)$
- $E_{MAP} = -(|T| + a - 1)\ln\theta - (|H| + b - 1)\ln(1-\theta)$
- $\theta_{MAP} = \frac{|T|+a-1}{|T|+|H|+a+b-2}$
- When $a = b = 1$: $\theta_{MAP} = \theta_{MLE}$

Posterior: $P(\theta|D) = \text{Beta}(\theta|a + |T|, b + |H|)$

Hoeffding's Inequality

$p(|\theta_{MLE} - \theta_{\text{true}}| \geq \epsilon) \leq 2e^{-2N\epsilon^2} \leq \delta$

Bayesian Models

Predictive Distribution: $p(f|D, a, b) = \int_0^1 p(f|\theta)p(\theta|D, a, b)d\theta$

Fully Bayesian: $\theta^* = \frac{|T|+a}{|T|+|H|+a+b} = \text{Ber}(f|\theta)$

Conjugate Priors

Bernoulli ⇌ Beta ($\Gamma(n) = (n-1)!$):
- Likelihood: $p(D|\theta) = \theta^k(1-\theta)^{n-k}$
- Prior: $\text{Beta}(\theta|a,b) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)}\theta^{a-1}(1-\theta)^{b-1}$
- When to use Bernoulli: Modeling binary outcome probabilities (success/failure, yes/no, click/no-click); Properties: Support $\theta \in [0,1]$. Parameters $a, b > 0$ (shape). Mean $\mathbb{E}[\theta] = \frac{a}{a+b}$, mode $\frac{a-1}{a+b-2}$ (for $a, b > 1$), variance $\frac{ab}{(a+b)^2(a+b+1)}$
- When to use Beta: Modeling any probability/proportion (not necessarily from Bernoulli data). Examples: click-through rates, conversion rates, exam pass rates, market share percentages

Poisson ⇌ Gamma:
- Likelihood: $p(D|\lambda) = \prod_i^N \frac{\lambda^{x_i} e^{-\lambda}}{x_i!}$ (a Poisson distribution _becomes_ Gaussian when the mean is large)
- Prior: $p(\lambda) = \frac{\beta^\alpha}{\Gamma(\alpha)}\lambda^{\alpha-1}e^{-\beta\lambda}$
- Posterior: $p(\lambda|D) = \text{Gamma}(\lambda|\alpha + \sum x_i, \beta + N)$
- When to use: Modeling count data or event rates (arrivals per hour, defects per unit, events per time interval)
- Properties: Support $\lambda \in (0, \infty)$. Parameters $\alpha > 0$ (shape), $\beta > 0$ (rate). Mean $\mathbb{E}[\lambda] = \frac{\alpha}{\beta}$, mode $\frac{\alpha-1}{\beta}$ (for $\alpha \geq 1$), variance $\frac{\alpha}{\beta^2}$.
- When to use Gamma: Modeling any positive continuous variable (not necessarily from Poisson data). Examples: time between events, product lifetimes, claim sizes, service times.

Gaussian ⇌ Gaussian:
- Likelihood: $p(D|\mu) = \prod_i^N \frac{1}{\sqrt{2\pi\sigma^2}}\exp\left(-\frac{(x_i-\mu)^2}{2\sigma^2}\right)$
- Prior: $p(\mu) = \mathcal{N}(\mu|\mu_0, \tau_0^2)$
- Posterior parameters:
- $\mu_{\text{post}} = (\mu_0\sigma^2 + N\bar{x}\tau_0^2)/(\sigma^2 + N\tau_0^2)$
- $\tau^2_{\text{post}} = 1/\left(\frac{1}{\tau_0^2} + \frac{N}{\sigma^2}\right)$
- When to use: Modeling continuous measurements with known variance (sensor readings, heights, temperatures, test scores)
- Properties: Support $\mu \in (-\infty, \infty)$. Parameters μ_0 (prior mean), τ_0^2 (prior variance), σ^2 (known data variance). Mean $\mathbb{E}[\mu] = \mu_{\text{post}}$, mode μ_{post} (symmetric), variance τ^2_{post}.

Precision form $\tau_{\text{post}}^{-2} = \tau_0^{-2} + N\sigma^{-2}$ (precisions add).

Uniform ⇌ Pareto:
- Likelihood: $p(D|\theta) = \theta^{-N} \cdot \mathbb{1}_{\max(x_i) \leq \theta}$
- Prior: $\text{Pareto}(\theta|\lambda, \alpha) = \frac{\alpha\lambda^\alpha}{\theta^{\alpha+1}} \cdot \mathbb{1}_{\theta \geq \lambda}$
- When to use Uniform: Modeling data with a hard, unknown upper bound (where data is equally likely anywhere below the bound). Examples: Serial number analysis (German Tank Problem), estimating maximum physical limits; Properties: Support $\theta \in [\lambda, \infty)$. Parameters $\lambda, \alpha > 0$. Mean $\mathbb{E}[\theta] = \frac{\alpha\lambda}{\alpha-1}$ (for $\alpha > 1$), mode λ, variance $\frac{\alpha\lambda^2}{(\alpha-1)^2(\alpha-2)}$ (for $\alpha > 2$).
- When to use Pareto: Modeling heavy-tailed quantities or the distribution of a minimum/maximum threshold. Examples: Wealth distribution, city populations, or (as here) the belief about the maximum possible value of a Uniform variable.

4. Linear Regression

Model: $y_i = f(x_i) + \varepsilon_i \quad \varepsilon_i \sim \mathcal{N}(0, \beta^{-1})$

Least Squares Error: $E_{LS} = \frac{1}{2}\sum_i^n (w^T x_i - y_i)^2$

Optimal Weight: $w^* = \arg\min_w E_{LS} = (X^T X)^{-1} X^T y = X^\dagger y$

Non-linear Data (Feature Transform):
- $f(x) = w_0 + \sum_{j=1}^M w_j \phi_j(x) = w^T \Phi(x)$
- $\Phi \in \mathbb{R}^{N \times (M+1)}$
- $w^* = (\Phi^T\Phi)^{-1}\Phi^T y = \Phi^\dagger y$

Model Complexity:
- High variance → overfit
- High bias → underfit

Ridge Regression: $E_{ridge} = \frac{1}{2}\sum_i^n (w^T\Phi(x_i) - y_i)^2 + \frac{\lambda}{2}\|w\|_2^2$

Probabilistic Formulation

Likelihood: $y_i \sim \mathcal{N}(f_w(x_i), \beta^{-1})$, $p(y|X, w, \beta) = \prod_i^N p(y_i|f_w(x_i), \beta)$

Negative Log-Likelihood:
- $E_{ML} = -\ln p(y|X, w, \beta)$
- $E_{ML} = \frac{\beta}{2}\sum_i^N (w^T\Phi(x_i) - y_i)^2 - \frac{N}{2}\ln\beta + \frac{N}{2}\ln 2\pi$

Maximum Likelihood Estimators:
- $w_{ML} = w_{LS} = \Phi^\dagger y$
- $\frac{1}{\beta_{ML}} = \frac{1}{N}\sum_i^N (w_{ML}^T \Phi(x_i) - y_i)^2$

With Gaussian Prior:

$p(w|\alpha) = \mathcal{N}(w|0, \alpha^{-1}I) = \left(\frac{\alpha}{2\pi}\right)^{\frac{M}{2}} \exp\left(-\frac{\alpha}{2}w^T w\right)$ (M: length of w)

MAP Estimation:
- $E_{MAP} = -\ln p(y|X, w, \beta) - \ln p(w|\alpha)$
- $E_{MAP} = \frac{\beta}{2}\sum_i^N (w^T\phi(x_i) - y_i)^2 + \frac{\alpha}{2}\|w\|_2^2$
- Equivalent to Ridge Regression where $\lambda = \frac{\alpha}{\beta}$
- $w^*_{ridge} = (\Phi^T\Phi + \lambda I)^{-1}\Phi^T y$

Fully Bayesian Linear Regression

Posterior: $p(w|D) = \mathcal{N}(w|\mu, \Sigma)$
- $\mu = \beta\Sigma\Phi^T y$
- $\Sigma^{-1} = \alpha I + \beta\Phi^T\Phi$

Predictions:
- MLE: $p(\hat{y}_{new}|x_{new}, w_{ML}, \beta_{ML}) = \mathcal{N}(\hat{y}_{new}|w_{ML}^T\phi(x_{new}), \beta_{ML}^{-1})$
- MAP: $p(\hat{y}_{new}|x_{new}, w_{MAP}, \beta) = \mathcal{N}(\hat{y}_{new}|w_{MAP}^T\phi(x_{new}), \beta^{-1})$
- Fully Bayesian:

$p(\hat{y}_{new}|x_{new}, D) = \mathcal{N}(\hat{y}_{new}|\mu^T\Phi(x_{new}), \beta^{-1} + \phi(x_{new})^T\Sigma\phi(x_{new}))$

Weighted Linear Regression

Objective (with weight r_i): $E_{weighted} = \frac{1}{2}\sum_i^N r_i(w^T\phi(x_i) - y_i)^2$

Optimal Weight: $w^*_{weighted} = (\Phi^T R\Phi)^{-1}\Phi^T R y$

5. Linear Classification

Zero-one Loss: $l_{01}(y, \hat{y}) = \sum_i^N \mathbb{I}(\hat{y}_i \neq y_i)$ (loss for incorrect predictions is 1)

Hyperplane: $f(x) = w^T x + w_0$
- Direction: w
- Distance from origin: $-\frac{w_0}{\|w\|}$
- Distance to Plane: The distance from the point x to the decision boundary: $\frac{y(x)}{\|w\|}$

Perceptron Update Rule (for each misclassified x_i):

$w \leftarrow \begin{cases} w + x_i & \text{if } y_i = 1 \\ w - x_i & \text{if } y_i = 0 \end{cases} \quad w_0 \leftarrow \begin{cases} w_0 + 1 & \text{if } y_i = 1 \\ w_0 - 1 & \text{if } y_i = 0 \end{cases}$

Probabilistic Generative Model:
- Prior: $y \sim \text{Categorical}(\theta)$, $p(y = c) = \theta_c = \frac{N_c}{N}$, $\sum_c \theta_c = 1$
- Class-conditional: $p(x|y = c) = \mathcal{N}(x|\mu_c, \Sigma)$ (assume Σ_c all equal)

Probabilistic Generative Models & Discriminant Analysis

- Remember $\sum_{c=1}^C \sum_{n=1}^N = 1$

Binary Classification:
- $p(y = 1|x) = \sigma(a) = \frac{1}{1+e^{-a}}$ where $a = \ln\frac{p(x|y=1)p(y=1)}{p(x|y=0)p(y=0)}$
- $a = w^T x + w_0$

LDA (Linear Discriminant Analysis) (with shared covariance Σ):
- $w = \Sigma^{-1}(\mu_1 - \mu_0)$
- $w_0 = -\frac{1}{2}\mu_1^T\Sigma^{-1}\mu_1 + \frac{1}{2}\mu_0^T\Sigma^{-1}\mu_0 + \ln\frac{p(y=1)}{p(y=0)}$
- Thus $y|x \sim \text{Bernoulli}(\sigma(w^T x + w_0))$

Naive Bayes (assumes feature independence):
- $p(x|y = c) = \prod_{j=1}^d p(x_j|y = c)$
- Gaussian Naive Bayes: $p(x_j|y = c) = \mathcal{N}(x_j|\mu_{jc}, \sigma_{jc}^2)$
- Equivalent to LDA/QDA with diagonal covariance matrix
- $a = w^T x + w_0$ where $w_j = \frac{\mu_{j1}}{\sigma_{j1}^2} - \frac{\mu_{j0}}{\sigma_{j0}^2}$ (if $\sigma_{j1} = \sigma_{j0}$)
- Multinomial Naive Bayes: for discrete features (e.g., word counts)
- $p(x_j|y = c) = \text{Categorical}(\theta_{jc})$
- $\log p(y = c|x) \propto \sum_j x_j \log\theta_{jc} + \log p(y = c)$
- Bernoulli Naive Bayes: for binary features
- $p(x_j|y = c) = \text{Bernoulli}(\theta_{jc})$

Multi-class Classification:
- $p(y = c|x) = \frac{p(x|y=c)p(y=c)}{\sum_j^C p(x|y=c_j)p(y=c_j)} = \frac{\exp(w_c^T x + w_{c0})}{\sum_j^C \exp(w_j^T x + w_{j0})}$
- $w_c = \Sigma^{-1}\mu_c$
- $w_{c0} = -\frac{1}{2}\mu_c^T\Sigma^{-1}\mu_c + \ln p(y = c)$

QDA (Quadratic Discriminant Analysis) (different covariances Σ_c):
- $p(y = 1|x) = \sigma(a)$ where $a = x^T W_2 x + w_1^T x + w_0$
- $W_2 = \frac{1}{2}[\Sigma_0^{-1} - \Sigma_1^{-1}]$
- $w_1 = \Sigma_1^{-1}\mu_1 - \Sigma_0^{-1}\mu_0$
- $w_0 = -\frac{1}{2}\mu_1^T\Sigma_1^{-1}\mu_1 + \frac{1}{2}\mu_0^T\Sigma_0^{-1}\mu_0 + \ln\frac{\pi_1}{\pi_0} + \frac{1}{2}\ln\frac{|\Sigma_0|}{|\Sigma_1|}$

one hot
- $\prod_{c=1}^C (p_c)^{y_c} = p_{\text{true class}}$
- $p(\mathcal{D} \mid \pi_c, \theta_{c\,c=1}^{\,C}) = \prod_{n=1}^N \prod_{c=1}^C \left[p(x^{(n)}|\theta_c)\pi_c\right]^{y_c^{(n)}}$
- $\pi_c, \theta_{c\,c=1}^{\,C}$ = all parameters for all C classes

Linear Discriminant Model: Logistic Regression

Binary Logistic Regression:
- $p(y = 1|x) = \sigma(w^T x)$
- $p(y = 0|x) = 1 - \sigma(w^T x)$
- $p(y|w, x) = \prod_i^N \sigma(w^T x_i)^{y_i}(1 - \sigma(w^T x_i))^{1-y_i}$
- $\frac{d\sigma(a)}{da} = \sigma(a)(1 - \sigma(a))$

Loss Function (Binary Cross Entropy):
- $E(w) = -\sum_i^N (y_i \log\sigma(w^T x_i) + (1 - y_i)\log(1 - \sigma(w^T x_i)))$
- Regularization can be added: $+\lambda\|w\|_2^2$

Multi-class (Softmax + Cross Entropy):
- $E(w) = -\sum_i^N \sum_c^C y_{ic} \log\frac{e^{(w_c^T x)}}{\sum_{c'} e^{(w_{c'}^T x)}}$
- $y_{ic} = 1$ iff sample $x \in c$ class

6. Optimization

Convexity

A function is convex if:

1. $f((1-t)x + ty) \leq (1-t)f(x) + tf(y)$ (any point between two points is lower than the line connecting them)
2. $f(y) - f(x) \geq \frac{f((1-t)x+ty) - f(x)}{t}$
3. $f(y) \geq f(x) + (y - x)^T\nabla f(x)$
4. Hessian Matrix is positive semi-definite:

$$\nabla^2 f(x) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix} \succeq 0$$

(i.e., $v^T\nabla^2 f(x)v \geq 0$ for all vectors v)

1. Sylvester's Criterion for PSD (Convexity Check)

For Hessian, function is convex ⟺ BOTH 1: $A \geq 0$ 2: $AC - B^2 \geq 0$

(If det = 0, condition 1 is essential (prevents saddle points))

2. If det ≤ 0 not PSD

Properties of Convex Functions

Properties of Convex Functions:

1. $g(x) = g_1(x) + g_2(x)$ (Sum of convex functions is convex)
2. $g(x) = \max\{g_1(x), g_2(x)\}$ (Pointwise maximum is convex)
3. $g(x) = c \cdot g_1(x)$ where $c \geq 0$ (Non-negative scaling)
4. $g(x) = c \cdot f_1(x)$ where $c \leq 0$ and $f_1(x)$ is concave
5. $g(x) = g_1(Ax + b)$ (Affine transformation)

Gradient Descent (Line Search)

1. $\Delta\theta = -\nabla f(\theta)$
2. $t^* = \arg\min_{t \geq 0} f(\theta + t\Delta\theta)$
3. $\theta = \theta + t^*\Delta\theta$

SGD (Stochastic Gradient Descent)

$\theta = \theta - r \cdot \nabla f(\theta)$ where r is learning rate

Decaying learning rate: $r = \alpha r, \quad 0 < \alpha < 1$

Momentum

- $m_t = r \cdot \nabla f(\theta_t) + \gamma \cdot m_{t-1}$
- $\theta_{t+1} = \theta_t - m_t$

Adam

- $m_t = \beta_1 m_{t-1} + (1 - \beta_1)\nabla f(\theta_t)$ (mean)
- $v_t = \beta_2 v_{t-1} + (1 - \beta_2)(\nabla f(\theta_t))^2$ (variance)
- $\hat{m}_t = \frac{m_t}{1-\beta_1^t} \quad \hat{v}_t = \frac{v_t}{1-\beta_2^t}$
- $\theta_{t+1} = \theta_t - \frac{r}{\sqrt{\hat{v}_t + \varepsilon}}\hat{m}_t$
- Default values: $\beta_1 = 0.9, \beta_2 = 0.999, \varepsilon = 10^{-8}$

Newton Method

Taylor expansion: $f(\theta_t + \delta) = f(\theta_t) + \delta^T\nabla f(\theta_t) + \frac{1}{2}\delta^T\nabla^2 f(\theta_t)\delta + \ldots$

Update: $\theta_{t+1} = \theta_t - [\nabla^2 f(\theta_t)]^{-1}\nabla f(\theta_t)$

Mini-batch SGD

- $\theta_{t+1} = \theta_t - r \cdot \frac{n}{|S|}\sum_{j \in S}\nabla L_j(\theta_t)$
- Batch size ↓ ⟶ variance ↑
- Batch size ↑ ⟶ computation time ↑

7. Deep Learning

Notation: w_{ijk} denotes weight from layer i, input node j, output node k

Architecture Types

- Feed-Forward Neural Network (FFNN)
- Multi-layered Perceptron (MLP)

Activation Functions

- Sigmoid: $\sigma(x) = \frac{1}{1+e^{-x}}$
- ReLU: $\max(0, x)$
- ELU: $\begin{cases} x & x > 0 \\ a(e^x - 1) & x < 0 \end{cases}$
- tanh: $\tanh(x)$
- Leaky ReLU: $\max(0.1x, x)$
- Swish: $x \cdot \sigma(x)$

Target and Loss Functions

Target $p(y\|x)$	Activation	Loss	Formula
Binary (Bernoulli)	Sigmoid	BCE	$-[y\log\hat{y} + (1-y)\log(1-$
Discrete (Categorical)	Softmax	CE	$-\sum i\, y_i \log\hat{y}_i$
Continuous (Gaussian)	Identity	MSE	$\frac{1}{2}(y - \hat{y})^2$

Unified gradient: $\frac{\partial\mathcal{L}}{\partial z} = \hat{y} - y$

Weight Update Rule

$W^{(new)} = W^{(old)} - r\nabla_W E(W^{(old)})$

Backpropagation

Chain Rule: $\frac{\partial c}{\partial x} = \frac{\partial c}{\partial a}\frac{\partial a}{\partial x} + \frac{\partial c}{\partial b}\frac{\partial b}{\partial x}$

Gradient of vector:

$\nabla_a c = \left(\frac{\partial c}{\partial a}\right)^T = \left[\frac{\partial c}{\partial a_1}, \frac{\partial c}{\partial a_2}, \ldots, \frac{\partial c}{\partial a_m}\right]^T \in \mathbb{R}^{1 \times m}$

Vector chain rule: $\nabla_x c = \left(\frac{\partial a}{\partial x}\right)^T\nabla_a c$

Derivative Dimensions (Output vs. Input):

	scalar	vector	matrix
scalar	scalar	vector	matrix
vector	vector	matrix	3-way tensor
matrix	matrix	3-way tensor	4-way tensor

Tensor notation: $\left(\frac{\partial a}{\partial W}\right)_{ijk} = \frac{\partial a_i}{\partial W_{jk}}$

Coding

- Dot Product: x @ y (vectors ⟶ scalar, matrices ⟶ matrix)
- Element-wise Mul: x * y (vectors ⟶ vector)
- Transpose: x.T or W.T (crucial for matrix calculus)
- Safe Log: np.log1p(x) safer than np.log(1+x)
- Summation: np.sum(x, axis=0) (backprop from scalar loss to vector weights)

8. CNN & Deep Learning Architecture

CNN Kernel Parameters: $L \times M \times C_{in} \times C_{out}$

Padding

- VALID: No padding, $D_{l+1} = (D_l - K) + 1$
- SAME: Add $P = \lfloor\frac{K}{2}\rfloor$ padding on each side to keep size
- FULL: Add $P = K - 1$ on each side to increase size

Stride & Pooling

Stride: Step size (> 1 results in downsampling)
- $D_{l+1} = \lfloor\frac{D_l + 2P - K}{S}\rfloor + 1$

Pooling:
- Max pooling: Take maximum value
- Mean pooling: Take average value

Initialization & Training Issues

Gradient Problems:
- Vanishing gradient: W becomes too small
- Exploding gradient: W becomes too large

Xavier Initialization:
- $\text{Var}(W) = \frac{2}{fan\_in + fan\_out}$
- Uniform: $W \sim \text{Uniform}\left(-\sqrt{\frac{6}{fan\_in+fan\_out}}, \sqrt{\frac{6}{fan\_in+fan\_out}}\right)$
- Normal: $W \sim \mathcal{N}\left(0, \frac{2}{fan\_in+fan\_out}\right)$
- Used for saturating activations like sigmoid and tanh

Regularization & Normalization

- Regularization techniques:
- Adding L_2 norm (Weight Decay).
- Early stopping.
- Data augmentation.
- Injecting noise.
- Dropout: Used only during training.

Batch Normalization

- Standardizes inputs to a layer for each mini-batch:

$\hat{x} = \frac{x - E_B(x)}{\sqrt{\text{Var}_B(x) + \epsilon}} \iff x = \gamma\hat{x} + \beta$

Residual Learning (Skip Connections)

- Skip Connection formula: $y = f(x, W)T(x, W) + x(1 - T(x, W))$
- This allows gradients to flow through the network more easily, facilitating the training of very deep networks.

9. Support Vector Machines (SVM)

Margin: $\frac{2}{\|w\|}$

Constraints:

- $w^T x_i + b \geq 1$ for $y_i = +1$
- $w^T x_i + b \leq -1$ for $y_i = -1$
- Thus: $y_i(w^T x_i + b) - 1 \geq 0$ for $\forall x_i$

Optimization Problem:

- Minimize: $\frac{1}{2} w^T w$
- Subject to: $f_i(w, b) = y_i(w^T x_i + b) - 1 \geq 0$

Lagrangian Dual Function

- Dual function: $g(\alpha) = \min_{\theta \in \mathbb{R}^d} \left(f_0(\theta) + \sum_{i=1}^M \alpha_i f_i(\theta) \right)$
- Lagrangian: $L(\theta, \alpha) = f_0(\theta) + \sum_{i=1}^M \alpha_i f_i(\theta)$
- Conditions: $\alpha_i \geq 0$ and $f_i(\theta) \leq 0$

SVM Optimization Steps

1. Calculate Lagrangian:

- $L(w, b, \alpha) = \frac{1}{2} w^T w - \sum_i^N \alpha_i [y_i(w^T x_i + b) - 1]$

2. Minimize L:

- $\frac{\partial L}{\partial w} = w - \sum_i^N \alpha_i y_i x_i \overset{!}{=} 0$
- $\frac{\partial L}{\partial b} = \sum_i^N \alpha_i y_i \overset{!}{=} 0$
- $\Rightarrow w = \sum_i^N \alpha_i y_i x_i$

3. Dual Problem:

- $g(\alpha) = \sum_i^N \alpha_i - \frac{1}{2} \sum_i^N \sum_j^N y_i y_j \alpha_i \alpha_j x_i^T x_j$
- Note: $x_i^T x_j$ can be replaced by Kernel $\Phi(x_i, x_j)$
- w.r.t. $\alpha_i \geq 0, \quad \sum_i^N \alpha_i y_i = 0$
- $w = \sum_i^N \alpha_i^* y_i x_i$
- $b = \frac{1}{y_i} - w^T x_i = y_i - w^T x_i$
- $\therefore h(x) = \operatorname{sign}(\sum_{i \in S} \alpha_i y_i x_i^T x + b)$

Soft SVM (relaxed margin):

- Constraint: $y_i(w^T x_i + b) \geq 1 - \xi_i \quad (\xi_i \geq 0)$
- Minimize: $f_0(w, b, \xi) = \frac{1}{2} w^T w + C \sum_i^N \xi_i$
- w.r.t.:
- $y_i(w^T x_i + b) - 1 + \xi_i \geq 0$
- $\xi_i \geq 0$
- ($C \to \infty$: hard margin)

Soft Margin SVM Derivation

1. Lagrangian Formulation:

$$L(w, b, \xi, \alpha, \mu) = \frac{1}{2} w^T w + C \sum_i^N \xi_i - \sum_i^N \alpha_i [y_i(w^T x_i + b) - 1 + \xi_i] - \sum_i^N \mu_i$$

2. Minimize L (Gradients):

- $\frac{\partial L}{\partial w} = w - \sum_i^N \alpha_i y_i x_i \overset{!}{=} 0$
- $\frac{\partial L}{\partial b} = \sum_i^N \alpha_i y_i \overset{!}{=} 0$
- $\frac{\partial L}{\partial \xi_i} = C - \alpha_i - \mu_i \overset{!}{=} 0 \Rightarrow \alpha_i = C - \mu_i$
- Since $\mu_i \geq 0$ and $\alpha_i \geq 0$: Box Constraint $\alpha_i \in [0, C]$

3. Maximize Dual Function:

$$g(\alpha) = \sum_i^N \alpha_i - \frac{1}{2} \sum_i^N \sum_j^N y_i y_j \alpha_i \alpha_j x_i^T x_j$$

Subject to:

- $\sum_i^N \alpha_i y_i = 0$
- $0 \leq \alpha_i \leq C$

Interpretation of α_i:

- If $0 < \alpha_i < C$: $\xi_i = 0$ and $y_i(w^\top x_i + b) = 1$ (point lies exactly on the margin)
- If $\alpha_i = C$: $\xi_i > 0$ (point violates the margin)
- If $0 < \xi_i < 1$: $0 < y_i(w^\top x_i + b) < 1$, point is inside the margin but correctly classified
- If $\xi_i \geq 1$: $y_i(w^\top x_i + b) \leq 0$, point is misclassified
- Larger $C \longrightarrow$ less tolerance for points inside the margin

Hinge Loss:

$$\frac{1}{2} w^T w + C \sum_i^N \max\{0, 1 - y_i(w^T x_i + b)\}$$

Kernel Methods

Definition: $k(x_i, x_j) = \phi(x_i)^T \phi(x_j)$

Prediction: $h(x) = \operatorname{sign}\left(\sum_{j \in S} \alpha_i y_i k(x_i, x) + b\right)$

Kernel Matrix Properties, Mercer: Must be symmetric positive semi-definite

$c^T K c \geq 0$

$\text{Sum} = \sum_i \sum_j c_i c_j k(\mathbf{x_i}, \mathbf{x_j}) \geq 0$

Valid Kernel Construction Rules:

1. Sum: $k(x_1, x_2) = k_1 + k_2$
2. Scaling: $k(x_1, x_2) = c k_1$ with $c > 0$
3. Product: $k(x_1, x_2) = k_1 \cdot k_2$
4. Transformation: $k(x_1, x_2) = k_3(\phi(x_1), \phi(x_2))$
5. Matrix Scaling: $k(x_1, x_2) = x_1^T A x_2$ where A is symmetric positive-definite

Common Examples:

- Polynomial: $k(a, b) = (a^T b)^n$ or $(a^T b + 1)^P$
- Gaussian (RBF): $k(a, b) = \exp\left(-\frac{\|a - b\|^2}{2\sigma^2}\right)$

Mercer Theorem: $\text{Sum} = \sum_i \sum_j c_i c_j k(\mathbf{x}_i, \mathbf{x}_j) \geq 0$ / The Gram matrix must be PSD" ($c^T K c \geq 0$) see Convexity Section

Multiclass Classification Strategies

- 1 vs n classification: Look at maximum distance
- 1 vs 1 classification: Look at majority vote

10. Dimension Reduction PCA SVD

Dimension Reduction (PCA)

Transformation: $\tilde{x}' = \tilde{x} \cdot F \quad \Sigma_{\tilde{x}'} = F^T \Sigma_{\tilde{x}} F$

(Minimizing Error = Maximizing Variance)

1. Centering Data

$$\tilde{x}_i = x_i - \bar{x} \text{ where } \bar{x} = \begin{bmatrix} \bar{x}_1 \\ \vdots \\ \bar{x}_d \end{bmatrix} = \frac{1}{N} \cdot X^T \cdot 1_N$$

2. Variance and Covariance

- Variance: $\operatorname{Var}(X_j) = \frac{1}{N} X_j^T X_j - \bar{x}_j^2$
- Covariance: $\operatorname{Cov}(X_i, X_j) = \frac{1}{N} X_i^T X_j - \bar{x}_i \bar{x}_j$
- Covariance Matrix: $\Sigma_{\tilde{X}} = \frac{1}{N} \tilde{X}^T \tilde{X}$ (symmetric)

3. Eigen-decomposition

$\Sigma_{\tilde{X}} = \Gamma \Lambda \Gamma^T$ (where Λ is diagonal)

4. Transformation

- $Y = \tilde{X} \cdot \Gamma$
- $Y_{reduced} = \tilde{X} \cdot \Gamma_{truncated}$
- Reconstruction: $\tilde{X}_{reconstructed} = Y_{reduced} \cdot \Gamma_{truncated}^T$ (back to original dims)
- Criteria: Retain 90% of variance: $\sum_{i=1}^k \lambda_i \geq 0.9 \sum_{i=1}^d \lambda_i$
- Complexity: $O(nd^2 + d^3)$

5. Iterative Eigenvector Calculation

Power Iteration: $v \leftarrow \frac{A \cdot v}{\|A \cdot v\|}$ (converges to eigenvector with largest eigenvalue)

Singular Value Decomposition (SVD)

Goal: Find best low-rank approximation of matrix A

Frobenius Norm Objective: $\|A - B\|_F^2 = \sum_i^n \sum_j^D (a_{ij} - b_{ij})^2$

Complexity: $O(n \cdot d^2)$ or $O(n^2 \cdot d)$

Decomposition: $A = U \Sigma V^T$ where:

- $U \in \mathbb{R}^{n \times r}$ (user-to-concept similarity)
- $\Sigma \in \mathbb{R}^{r \times r}$
- $V \in \mathbb{R}^{d \times r}$ (item-to-concept similarity)

Rank-2 Decomposition: $A = (\sigma_1 \cdot u_1 \cdot v_1^T) + (\sigma_2 \cdot u_2 \cdot v_2^T)$

EYM

$\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T = \sum_{i=1}^{\min(m,n)} \sigma_i \mathbf{u}_i \mathbf{v}_i^T$

Using SVD for Dimensionality Reduction

Projection: $P = U\Sigma$ or $P = A \cdot V$

Reconstruction: $A_{reconstructed} = U_k \Sigma_k V_k^T$ or

$A_{reconstructed} = P \cdot V_k^T$

Retain 90% energy: $\sum_{i=1}^k \sigma_i^2 \geq 0.9 \sum_{i=1}^r \sigma_i^2$

Relationship to Eigenvectors:

- V contains eigenvectors of $X^T X$
- U contains eigenvectors of $X X^T$

SVD vs PCA

- Eigenvectors = Singular Vectors; (N)Eigenvalues = (Singular Values)^2
- Transform the data such that dimensions of new space are uncorrelated + discard (new) dimensions with smallest variance = find optimal low-rank approximation(norm_F)

Matrix Factorization (MF)

1. Fundamentals & Metrics

RMSE: $\sqrt{\frac{1}{|S|} \sum (r_{ui} - \hat{r}_{ui})^2}$

SSE: $\sum (r_{ui} - [U \Sigma V^T]_{ui})^2$

Decomposition: $R = U \Sigma V^T \approx Q \cdot P^T$

Prediction: $\hat{r}_{ui} = q_u \cdot p_i^T$

2. Alternating Optimization

1. Initialize: $P_0, Q_0, t = 0$
2. Update P: $P_{t+1} = \operatorname{argmin}_P f(P, Q_t)$

- Closed form: $p_i^T = \left(\frac{1}{|S_{*,i}|} \sum q_u^T q_u \right)^{-1} \cdot \frac{1}{|S_{*,i}|} \sum q_u^T r_{ui}$

3. Update Q: $Q_{t+1} = \operatorname{argmin}_Q f(P_{t+1}, Q)$
4. Repeat until convergence

3. Stochastic Gradient Descent (SGD)

- $e_{ui} \leftarrow r_{ui} - q_u \cdot p_i^T$
- $q_u \leftarrow q_u + 2r(e_{ui} p_i)$
- $p_i \leftarrow p_i + 2r(e_{ui} q_u)$ (r: learning rate)

4. Extensions: Bias & Regularization

Regularized Objective:

$\sum (r_{ui} - q_u \cdot p_i^T)^2 + \lambda_1 \sum \|q_u\|^2 + \lambda_2 \sum \|p_i\|^2$

Full Loss with Bias:

$L = \sum (r_{ui} - (q u p_i^T + bu + b_i + b))^2 + \lambda_1 \sum \|q u\|^2 + \lambda_2 \sum \|p_i\|^2$

SGD Updates (with Bias & Regularization):

1. $e_{ui} = r_{ui} - q_u \cdot p_i^T$
2. $q_u = q_u + 2r(e_{ui} p_i - \lambda_1 q_u)$
3. $p_i = p_i + 2r(e_{ui} q_u - \lambda_2 p_i)$
4. $b_u = b_u + 2re_{ui}$
5. $b_i = b_i + 2re_{ui}$
6. $b = \frac{1}{|S|} \sum r_{ui}$ (Global Bias)

11. Dimension Reduction Neighbor Graph Method

Neighbor Graph Method

Preserve local structure

Scatter Plot (High Dim) vs Adjacency Matrix (Graph)

t-SNE

High-dimensional similarity:

$$p_{j|i} = \frac{\exp(-\frac{\|x_i - x_j\|^2}{2\sigma_i^2})}{\sum_{k \neq i} \exp(-\frac{\|x_i - x_k\|^2}{2\sigma_i^2})}$$

- Data X
- $p_{ii} = 0$, $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2N}$
- $p_{ij} = p_{ji}$

Low-dimensional similarity:

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_k \sum_{l \neq k} (1 + \|y_k - y_l\|^2)^{-1}}$$

- Target y
- $q_{ii} = 0$

Next: minimize KL Divergence

$$\sum_i \sum_{j \neq i} p_{ij} \log \frac{p_{ij}}{q_{ij}} = KL(P||Q)$$

- $KL(P||Q) \geq 0$, $= 0$ iff $P = Q$

Divergence Behavior:

KL(P||Q) : Mean-seeking (Covers the whole distribution)

KL(P||Q) : Mode-seeking (Locks onto specific modes)

Autoencoder

Notes:

- Bottleneck (L << D)
- Extract key features and reconstruct; equivalent to only taking rank(X) = L.
- Formula: $X W_1 W_2 = \hat{X}$
- $W_1 W_2 = W^* = \Gamma$
- Γ is the largest L eigenvectors of $X^T X$.

12. Clustering

K-means: Objective (Distortion Measure)

- Formula:

$$J(X, Z, \mu) = \sum_i^N \sum_k^K z_{ik} \|x_i - \mu_k\|^2$$

- Indicator Variable:
- $z_i = \{0, 0, 1, 0, 0, 0\}$ implies sample i belongs to class 3.

Lloyd's Algorithm (Alternating Optimization)

1. Initialize: All centroids μ_i.
2. Assignment: Assign cluster indicators based on the nearest neighbor.

- _(Note: Whichever centroid a point is closest to, it belongs to that cluster.)_

3. Update:

- $\mu_k = \frac{1}{N_k} \sum_i^N z_{ik} x_i$
- Where $N_k = \sum_i^N z_{ik}$

4. Loop: Repeat until convergence.

K-means ++

1. Randomly select one data point as μ_1.
2. Calculate the distance of the remaining points to it: $\|x_i - \mu_1\|^2$.
3. Sample the next center point, with probability proportional to the distance size.
4. Re-calculate $D_i^2 = \min\{\|x_i - \mu_1\|^2, \|x_i - \mu_2\|^2, \ldots\}$.
5. Repeat steps 3 and 4 until K centers are chosen.

Gaussian Mixture Model (GMM)

Model Definition:

$p(x|\theta) = \sum_z p(x|z, \theta) p(z|\theta)$

- z is latent variables
- Optimization Goal: $\theta^* = \operatorname{argmax} \theta p(x|\theta)$

Distributions:

1. Prior: $p(z|\theta) = \operatorname{Cat}(\pi) \quad // \quad \{\pi_1, \pi_2, \ldots, \pi_K\}$

- Constraint: $\sum_i^K \pi_i = 1$

2. Likelihood: $p(x|z = k, \theta) = \mathcal{N}(x|\mu_k, \Sigma_k)$

Parameters:

Thus $\theta = \{\pi, \mu, \Sigma\}$

- π: k parameters (effectively $k - 1$)
- $\mu \in \mathbb{R}^d$: $k \cdot d$ parameters
- $\Sigma_k \in \mathbb{R}^{d \times d}$: $k \cdot d^2$ parameters (or $k \cdot \frac{d(d+1)}{2}$ due to symmetry)

Using GMM as a Generative Model

1. Sample class according to $\pi = (\pi_1, \pi_2, \ldots, \pi_K)$
2. Generate x according to $x \sim \mathcal{N}(\mu_k, \Sigma_k)$ (Probability Density)

Marginal Probability:

$$p(x|\pi, \mu, \Sigma) = \sum_i^K \pi_k \cdot \mathcal{N}(x|\mu_k, \Sigma_k)$$

Log-Likelihood:

$$\log p(X|\pi, \mu, \Sigma) = \sum_i^N \log \left(\sum_k^K \pi_k \cdot \mathcal{N}(x_i|\mu_k, \Sigma_k) \right)$$

Inference

Calculating the **responsibility** of component k for observation i:

$$p(z_{ik} = 1|x_i, \pi, \mu, \Sigma) = \frac{\pi_k \mathcal{N}(x_i|\mu_k, \Sigma_k)}{\sum_j^K \pi_j \mathcal{N}(x_i|\mu_j, \Sigma_j)}$$

Notation:

$\gamma(z_{ik}) = p(z_{ik} = 1|x_i, \pi, \mu, \Sigma)$

Expectation Maximization (EM) for GMM

Learning Process:

1. Initialize: $\pi^0, \mu_1^0 \ldots \mu_K^0, \Sigma_1^0 \ldots \Sigma_K^0$
2. E-step (Expectation):

Calculate the responsibility (posterior probability):

$$\gamma_t(z_{ik}) = \frac{\pi_k^t \mathcal{N}(x_i|\mu_k^t, \Sigma_k^t)}{\sum_j^K \pi_j^t \mathcal{N}(x_i|\mu_j^t, \Sigma_j^t)}$$

- Note: E-step evaluates the posterior $p(z|x, \theta^t)$.

3. M-step (Maximization):

Update parameters to maximize expected joint probability:

- $\mu_k^{t+1} = \frac{1}{N_k} \sum_i^N \gamma_t(z_{ik}) x_i$
- $\Sigma_k^{t+1} = \frac{1}{N_k} \sum_i^N \gamma_t(z_{ik})(x_i - \mu_k^{t+1})(x_i - \mu_k^{t+1})^T$
- $\pi_k^{t+1} = \frac{N_k}{N}$ where $N_k = \sum_i^N \gamma_t(z_{ik})$

4. Repeat steps 2 and 3 until convergence

Theoretical Objective:

- M-step Goal: $\theta^{t+1} = \operatorname{argmax}_\theta E_{z \sim p(z|x, \theta t)}[\log p(X, z|\theta)]$
- Decomposition: $\log p(X|\theta) = L(q, \theta) + KL(q||p(\cdot|X, \theta))$
- $\log p(X|\theta) \geq L(q, \theta)$ (Lower Bound)
- $L(q, \theta) = E_{z \sim q}[\log \frac{p(X, z|\theta)}{q(z)}] = \cdots + H[q]$

Model Selection & Hierarchical Clustering

Choosing K (Number of Clusters):

1. Bayesian Information Criterion (BIC)

- Formula: $BIC = M \log N - 2 \log \hat{L}$
- $M = K(1 + D + D^2)$ (Number of parameters)
- N: Sample size
- $\hat{L}$: Observed log-likelihood
- Goal: Minimize BIC.

2. AIC (Akaike Information Criterion)

- Formula: $AIC = 2M - 2 \log \hat{L}$

3. Agglomerative Clustering (Hierarchical)

- Method: Merge clusters recursively from bottom up
- _Description_: Merge classes gradually based on distance.

Misc

Gaussian

Multivariate Gaussian PDF

$$p(x|\mu, \Sigma) = \frac{1}{(2\pi)^{D/2} |\det(\Sigma)|^{1/2}} \exp\left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$

Isotropic Gaussian & Priors

Isotropic Gaussian:

(Refers to Gaussian distributions with covariance proportional to identity, i.e., spherical)

Priors:

- Laplace Prior: $\exp(-\frac{\|w\|}{\lambda})$
- Gaussian Prior: $\exp(-\frac{\|w\|^2}{2\sigma^2})$

Key Differences:

- **Laplace encourages sparsity** (results in sparse solutions, like L1 regularization).
- **L1 regularization** promotes sparsity, producing axis-aligned boundaries and effectively removing irrelevant dimensions by collapsing ellipsoid axes.
- **Gaussian results in solutions around the mean** (like L2 regularization).
- **L2 regularization** reduces variance and is more resistant to outliers, leading to smoother decision boundaries and less eccentric ellipses
- Laplace results in solutions around the **median**.

(i) Spherical/Isotropic: $\Sigma = \sigma I$

Matrix: $\begin{pmatrix} \sigma & 0 \\ 0 & \sigma \end{pmatrix}$ / **Constraints:** $A = D$, $B = C = 0$ / **Shape:** Circles

Boundary: Shared⟶Linear (fails concentric) / Separate⟶Circular (handles donuts)

Use: Round blobs

(iii) Diagonal: Σ = diag

Matrix: $\begin{pmatrix} \sigma x & 0 \\ 0 & \sigma y \end{pmatrix}$ / **Constraints:** $B = C = 0$, $A \neq D$ ok /

Shape: Axis-aligned ellipses

Boundary: Shared⟶Linear (fails concentric, = Naive Bayes) / Separate⟶Quadratic axis-aligned (handles concentric via tight-in-wide)

Use: Axis-aligned clusters OR concentric tight/loose

(ii) Full/Unconstrained

Matrix: $\begin{pmatrix} a & b \\ b & d \end{pmatrix}$ / **Constraints:** None / **Shape:** Tilted ellipses

Boundary: Shared⟶Linear (=LDA, fails concentric) / Separate⟶Quadratic (=QDA, handles all concentric/tilted)

Use: Diagonal/tilted clusters OR tilted concentric

Global Rule: Shared Cov = Linear (never concentric) / Separate Cov = Quadratic (solves concentric by enclosing)

Calculus Rules

$(\frac{u}{v})' = \frac{u'v - uv'}{v^2}$

$\frac{d^n}{dx^n} a^x = a^x (\ln(a))^n$

$\int a^x dx = \frac{a^x}{\ln(a)} + C$

Inequality

Jensen Ineq.:

$\mathbb{E}[f(x)] \geq f(\mathbb{E}[x])$

Or in discrete summation notation for many points $x_1, x_2, \ldots, x_n$ with weights p_i that sum to 1:

$\sum p_i f(x_i) \geq f\left(\sum p_i x_i\right)$

Triangle Inequality: $|a + b| \leq |a| + |b|$.

Statistics

- $\operatorname{Var}(X) = E(X^2) - [E(X)]^2$
- $\operatorname{Var}(aX) = a^2 \operatorname{Var}(X)$ (squared!)
- $\operatorname{Var}(X - Y) = \operatorname{Var}(X) + \operatorname{Var}(Y) - 2\operatorname{Cov}(X, Y)$ (still plus Var(Y)!)
- If independent: $\operatorname{Var}(X - Y) = \operatorname{Var}(X) + \operatorname{Var}(Y)$ (not minus!)
- $\operatorname{Cov}(X, Y) = E(XY) - E(X)E(Y)$
- $E(XY) = E(X)E(Y)$ only if independent

KKT Conditions

Problem: $\min x \; f(x)$ s.t. $g_i(x) \leq 0, \; h_j(x) = 0$

Lagrangian: $\mathcal{L}(x, \lambda, \nu) = f(x) + \sum_i \lambda_i g_i(x) + \sum_j \nu_j h_j(x)$

Five Conditions:

1. Stationarity: $\nabla_x \mathcal{L}(x^*, \lambda^*, \nu^*) = 0$
2. Primal Feasibility: $g_i(x^*) \leq 0, \; h_j(x^*) = 0$
3. Dual Feasibility: $\lambda_i \geq 0$
4. Complementary Slackness: $\lambda_i g_i(x^*) = 0$
5. Constraint Qualification: LICQ/Slater/MFCQ

Key Facts:

- Necessary: always (if CQ holds) / Sufficient: only if convex
- Active ($g_i(x^*) = 0$) ⟶ $\lambda_i \geq 0$ / Inactive ($g_i(x^*) < 0$) ⟶ $\lambda_i = 0$

Steps: (1) Write $\mathcal{L}$ (2) Solve $\nabla_x \mathcal{L} = 0$ (3) Check cases (4) Verify all conditions

Series

$$\frac{1}{1 - z} = \sum_{n=0}^{\infty} z^n = 1 + z + z^2 + z^3 + \ldots$$

$$e^z = \sum_{n=0}^{\infty} \frac{z^n}{n!} = 1 + z + \frac{z^2}{2!} + \frac{z^3}{3!} + \ldots$$

$$(a + b)^d = \sum_{k=0}^{d} \binom{d}{k} a^{d-k} b^k$$

$\binom{n}{k} = \frac{n!}{k!(n-k)!}$

Linear Algebra

PSD: $\lambda \geq 0$, $x^T A x \geq 0$, $\det(A) \geq 0$, symmetric

Invertible: $\lambda \neq 0$, $\det(A) \neq 0$, full rank, $Ax = b$ unique

PD (both): $\lambda > 0$, $x^T A x > 0$, invertible + PSD

- Trace Tricks: $\frac{\partial \operatorname{Tr}(AB)}{\partial A} = B^T$; $\operatorname{tr}(\mathbf{ABC}) = \operatorname{tr}(\mathbf{BCA}) = \operatorname{tr}(\mathbf{CAB})$; $\operatorname{tr}(\mathbf{A}) = \operatorname{tr}(\mathbf{A}^T)$
- Sandwich Rule: If you see $(\mathbf{x} - \mathbf{a})^T \mathbf{M} (\mathbf{x} - \mathbf{a})$ where $\mathbf{M}$ is symmetric, you can instantly rewrite it as: $\mathbf{x}^T \mathbf{M} \mathbf{x} - 2\mathbf{x}^T \mathbf{M} \mathbf{a} + \mathbf{a}^T \mathbf{M} \mathbf{a}$; Deriv is $2\mathbf{M}\mathbf{x}$
- In Linear Algebra, the **Determinant** of a matrix is the product of its eigenvalues: