

2. DT KNN

Decision Tree

Impurity Measures:

$i_E(t) = 1 - \max_c \pi_c$

Entropy (Shannon): $i_H(t) = -\sum_i c_i \pi_{ci} \log_2 \pi_{ci}$

Gini index: $i_G(t) = 1 - \sum_i c_i \pi_{ci}^2$

Greedy Optimization: Use $i_H(t)$ or $i_G(t)$ (not $i_E(t)$, which doesn't decrease impurity)

Information Gain:

$$p_L = \frac{N_{left}}{N}, p_R = \frac{N_{right}}{N}$$

$$\Delta i = i(t) - p_L \cdot i(t_L) - p_R \cdot i(t_R)$$

Stopping Conditions: $i(t) = 0 / d_{max} / N_{node} < t_n / \Delta i(s, t) < t_s$

LOOCV: Equivalent to N -fold cross-validation

KNN

Prediction: $\hat{y} = \arg \max_c \sum_{x_i \in N_k(x)} \mathbb{I}(y_i = c)$

Distance Metrics:

L1 (Manhattan): $d(x_1, x_2) = \sum_i |x_{1i} - x_{2i}|$

L2 (Euclidean): $d(x_1, x_2) = \sqrt{\sum_i (x_{1i} - x_{2i})^2}$

L ∞ : $d(x_1, x_2) = \max_i |x_{1i} - x_{2i}|$

Cosine Similarity: $\text{Sim}(x_1, x_2) = \frac{x_1^T x_2}{\|x_1\| \|x_2\|}$

Mahalanobis Distance: $\sqrt{(x_1 - x_2)^T \Sigma^{-1} (x_1 - x_2)}$ (Σ positive semi-definite, symmetric)

Weighted KNN (inverse distance weighting, closer points more important):

$$\hat{y} = \arg \max_c \frac{1}{k} \sum_{x_i \in N_k(x)} \frac{1}{d(x_i, x)} \mathbb{I}(y_i = c)$$

Hyperparameter Selection:

k small \rightarrow overfitting

k large \rightarrow underfitting

Use odd number to avoid ties

Scale Issue: Normalization $x_f = \frac{x_i - \mu_i}{\sigma_i}$ (or use weighted distance)

Confusion Matrix

Ground \ Predict	1	0
1	TP	FN
0	FP	TN

Metrics:

Precision: $\frac{TP}{TP+FP}$

Sensitivity/Recall: $\frac{TP}{TP+FN}$

Accuracy: $\frac{TP+TN}{TP+TN+FP+FN}$

F1 Score: $\frac{2 \cdot \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$

3. Prob Method

Probabilistic Inference

Maximum Likelihood Estimation (MLE):

$\theta_{MLE} = \arg \max_{\theta} p(D|\theta)$

$p(D|\theta) = \prod_i p(x_i|\theta)$

$E_{MLE} = -\ln p(D|\theta) = -\sum_i \ln p(x_i|\theta)$

$\theta_{MLE} = \frac{|T|}{|T|+|H|}$

Maximum A Posteriori (MAP):

$\theta_{MAP} = \arg \max_{\theta} p(\theta|D)$

$p(\theta|D) \propto p(D|\theta)p(\theta)$

$E_{MAP} = -(|T| + a - 1) \ln \theta - (|H| + b - 1) \ln(1 - \theta)$

$\theta_{MAP} = \frac{|T|+a-1}{|T|+|H|+a+b-2}$

When $a = b = 1$: $\theta_{MAP} = \theta_{MLE}$

Posterior: $P(\theta|D) = \text{Beta}(\theta|a + |T|, b + |H|)$

Hoeffding's Inequality

$$p(|\theta_{MLE} - \theta_{true}| \geq \epsilon) \leq 2e^{-2N\epsilon^2} \leq \delta$$

Bayesian Models

Predictive Distribution: $p(f|D, a, b) = \int_0^1 p(f|\theta)p(\theta|D, a, b)d\theta$

Fully Bayesian: $\theta^* = \frac{|T|+a}{|T|+|H|+a+b} \equiv \text{Ber}(\theta|f)$

Conjugate Priors

Bernoulli \Rightarrow Beta ($\Gamma(n) = (n-1)!$):

Likelihood: $p(D|\theta) = \theta^t (1-\theta)^{n-t}$

Prior: $\text{Beta}(\theta|a, b) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} \theta^{a-1} (1-\theta)^{b-1}$

 Poisson \Leftrightarrow Gamma:

Likelihood: $p(D|\lambda) = \prod_i \frac{\lambda^{x_i} e^{-\lambda}}{x_i!}$

Prior: $p(\lambda) = \frac{\alpha^a}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda}$

Posterior: $p(\lambda|D) = \text{Gamma}(\lambda|\alpha + \sum_i x_i, \beta + N)$

Gaussian = Gaussian:

Likelihood: $p(D|\mu) = \prod_i \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right)$

Prior: $p(\mu) = \mathcal{N}(\mu|\mu_0, \tau_0^2)$

Posterior parameters:

$$\mu_{post} = (\mu_0 \sigma^2 + N \bar{x}_0^2) / (\sigma^2 + N \tau_0^2)$$

$$\sigma_{post}^2 = 1 / \left(\frac{1}{\tau_0^2} + \frac{N}{\sigma^2} \right)$$

4. Linear Regression

Model: $y_i = f(x_i) + \varepsilon_i \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$

Least Squares Error: $E_{LS} = \frac{1}{2} \sum_i (w^T x_i - y_i)^2$

Optimal Weight: $w^* = \arg \min_w E_{LS} = (X^T X)^{-1} X^T y = X^{\dagger} y$

Non-linear Data (Feature Transform):

$f(x) = w_0 + \sum_{j=1}^M w_j \phi_j(x) = w^T \Phi(x)$

$\Phi \in \mathbb{R}^{N \times (M+1)}$

$w^* = (\Phi^T \Phi)^{-1} \Phi^T y = \Phi^{\dagger} y$

Model Complexity:

High variance \rightarrow overfit

High bias \rightarrow underfit

Ridge Regression: $E_{ridge} = \frac{1}{2} \sum_i (w^T \Phi(x_i) - y_i)^2 + \frac{\lambda}{2} \|w\|^2$

Probabilistic Formulation

Likelihood: $y_i \sim \mathcal{N}(f_w(x_i), \sigma^2)$, $p(y|X, w, \sigma^2) = \prod_i p(y_i|f_w(x_i), \sigma^2)$

Negative Log-Likelihood:

$E_{ML} = -\ln p(y|X, w, \sigma^2)$

$$E_{ML} = \frac{1}{2} \sum_i (w^T \Phi(x_i) - y_i)^2 + \frac{N}{2} \ln 2\pi$$

Maximum Likelihood Estimators:

$$w_{ML} = w_{LS} = \Phi^{\dagger} y$$

$\cdot \frac{1}{E_{ML}} = \frac{1}{N} \sum_i^N (w_{ML}^T \Phi(x_i) - y_i)^2$

With Gaussian Prior: $p(w|\alpha) = \mathcal{N}(w|0, \alpha^{-1} I) = \left(\frac{\alpha}{2\pi}\right)^{\frac{M}{2}} \exp\left(-\frac{\alpha}{2} w^T w\right)$ (M : length of w)

MAP Estimation:

$$E_{MAP} = -\ln p(y|X, w, \beta) - \ln p(w|\alpha)$$

$$E_{MAP} = \frac{\alpha}{2} \sum_i^N (w^T \Phi(x_i) - y_i)^2 + \frac{\beta}{2} \|w\|_2^2$$

Equivalent to Ridge Regression where $\lambda = \frac{\alpha}{\beta}$

$$w_{ridge}^* = (\Phi^T \Phi + \lambda I)^{-1} \Phi^T y$$

Fully Bayesian Linear Regression

Posterior: $p(w|D) = \mathcal{N}(w|\mu, \Sigma)$

$$\mu = \beta \Sigma \Phi^T y$$

$$\Sigma^{-1} = \alpha I + \beta \Phi^T \Phi$$

Predictions:

$$\text{MLE: } p(\hat{y}_{new}|x_{new}, w_{ML}, \beta_{ML}) = \mathcal{N}(\hat{y}_{new}|w_{ML}^T \Phi(x_{new}), \beta_{ML}^{-1})$$

$$\text{MAP: } p(\hat{y}_{new}|x_{new}, w_{MAP}, \beta) = \mathcal{N}(\hat{y}_{new}|w_{MAP}^T \Phi(x_{new}), \beta^{-1})$$

Fully Bayesian:

$$p(\hat{y}_{new}|x_{new}, D) = \mathcal{N}(\hat{y}_{new}|w^T \Phi(x_{new}), \beta^{-1} + \Phi(x_{new})^T \Sigma \Phi(x_{new}))$$

Weighted Linear Regression

Objective (with weight r_i): $E_{weighted} = \frac{1}{2} \sum_i r_i (w^T \Phi(x_i) - y_i)^2$

Optimal Weight: $w^*_{weighted} = (\Phi^T R \Phi)^{-1} \Phi^T R y$

5. Linear Classification

Zero-one Loss: $l_{01}(y, \hat{y}) = \sum_i \mathbb{I}(\hat{y}_i \neq y_i)$ (loss for incorrect predictions is 1)

Hyperplane: $f(x) = w^T x + w_0$

Direction: w

Distance from origin: $-\frac{w_0}{\|w\|}$

Perceptron Update Rule (for each misclassified x_i):

$$w \leftarrow \begin{cases} w + x_i & \text{if } y_i = 1 \\ w - x_i & \text{if } y_i = 0 \end{cases} \quad w_0 \leftarrow \begin{cases} w_0 + 1 & \text{if } y_i = 1 \\ w_0 - 1 & \text{if } y_i = 0 \end{cases}$$

Probabilistic Generative Model:

Prior: $y \sim \text{Categorical}(\theta)$, $p(y=c) = \theta_c = \frac{N_c}{N}$, $\sum_c \theta_c = 1$

Class-conditional: $p(x|y=c) = \mathcal{N}(\mu_c, \Sigma_c)$ (assume Σ_c all equal)

Probabilistic Generative Models & Discriminant Analysis

Binary Classification:

$$p(y=1|x) = \sigma(a) = \frac{1}{1+e^{-a}} \text{ where } a = \ln \frac{p(x|y=1)p(y=1)}{p(x|y=0)p(y=0)}$$

LDA (Linear Discriminant Analysis) (with shared covariance Σ):

$$w = \Sigma^{-1}(\mu_1 - \mu_0)$$

$$w_0 = -\frac{1}{2} w^T \Sigma^{-1} \mu_1 + \frac{1}{2} w^T \Sigma^{-1} \mu_0 + \ln \frac{p(y=1)}{p(y=0)}$$

Thus $y|x \sim \text{Bernoulli}(\sigma(w^T x + w_0))$

Multi-class Classification:

$$p(y=c|x) = \frac{p(x|y=c)p(y=c)}{\sum_j^C p(x|y=j)p(y=j)} = \frac{\exp(w_c^T x + w_{c0})}{\sum_j^C \exp(w_j^T x + w_{j0})}$$

$$w_c = \Sigma^{-1} \mu_c$$

$$w_{c0} = -\frac{1}{2} \mu_c^T \Sigma^{-1} \mu_c + \ln \sigma(w_c)$$

QDA (Quadratic Discriminant Analysis) (different covariances Σ_c):

$$p(y=1|x) = \sigma(a) \text{ where } a = x^T W x + w^T x + w_0$$

$$W = \frac{1}{2} [\Sigma_1^{-1} - \Sigma_0^{-1}]$$

$$w_1 = \Sigma_1^{-1} \mu_1 - \Sigma_0^{-1} \mu_0$$

$$w_0 = -\frac{1}{2} \mu_1^T \Sigma_1^{-1} \mu_1 + \frac{1}{2} \mu_0^T \Sigma_0^{-1} \mu_0 + \ln \frac{\pi_1}{\pi_0} + \frac{1}{2} \ln \frac{|\Sigma_1|}{|\Sigma_0|}$$

Linear Discriminant Model: Logistic Regression

Binary Logistic Regression:

$$p(y=1|x) = \sigma(w^T x)$$

$$p(y=0|x) = 1 - \sigma(w^T x)$$

$$p(y|w, x) = \prod_i^N \sigma(w^T x_i) y_i (1 - \sigma(w^T x_i))^{1-y_i}$$

Loss Function (Binary Cross Entropy):

$$E(w) = -\sum_i^N \sum_c^C y_{ic} \log \frac{e^{w_c^T x_i}}{\sum_c^C e^{w_c^T x_i}}$$

$y_{ic} = 1$ iff sample $x \in c$ class

6. Optimization

Convexity:

A function is convex if:

1. $f((1-t)x + ty) \leq (1-t)f(x) + tf(y)$ (any point between two points is lower than the line connecting them)

$$2. f(y) - f(x) \geq \frac{f((1-t)x+ty) - f(x)}{t}$$

$$3. f(y) \geq f(x) + (y - x)^T \nabla f(x)$$

4. Hessian Matrix is positive semi-definite

Gradient Descent (Line Search):

$$1. \Delta \theta = -\nabla f(\theta)$$

$$2. t^* = \arg \min_{t \geq 0} f(\theta + t\Delta \theta)$$

$$3. \theta = \theta + t^* \Delta \theta$$

SGD (Stochastic Gradient Descent)

$\theta = \theta - r \cdot \nabla f(\theta)$ where r is learning rate

Decaying learning rate: $r = \alpha r, \quad 0 < \alpha < 1$

Momentum:

$$m_t = r \cdot \nabla f(\theta_t) + \gamma \cdot m_{t-1}$$

$$\theta_{t+1} = \theta_t - m_t$$

Adam:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla f(\theta_t) \quad (\text{mean})$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla f(\theta_t))^2 \quad (\text{variance})$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\theta_{t+1} = \theta_t - \frac{r}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$$

Default values: $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$

Newton Method:

Taylor expansion: $f(\theta_t + \delta) = f(\theta_t) + \delta^T \nabla f(\theta_t) + \frac{1}{2} \delta^T \nabla^2 f(\theta_t) \delta + \dots$

Update: $\theta_{t+1} = \theta_t - [\nabla^2 f(\theta_t)]^{-1} \nabla f(\theta_t)$

Mini-batch SGD:

$$\theta_{t+1} = \theta_t - r \cdot \frac{n}{|S|} \sum_{j \in S} \nabla L_j(\theta_t)$$

Batch size $\downarrow \rightarrow$ variance \uparrow

Batch size $\uparrow \rightarrow$ computation time \uparrow

Feed-Forward Neural Network (FFNN)

Multi-layered Perceptron (MLP)

Activation Functions:

$$\text{Sigmoid: } \sigma(x) = \frac{1}{1+e^{-x}}$$

$$\text{ReLU: } \max(0, x)$$

$$\text{ELU: } \begin{cases} x & x > 0 \\ \alpha e^x - 1 & x < 0 \end{cases}$$

$$\text{tanh: } \tanh(x)$$

$$\text{Leaky ReLU: } \max(0.1x, x)$$

$$\text{Swish: } x \cdot \sigma(x)$$

Target and Loss Functions

Target $y(x)$	Final Layer	
---------------	-------------	--

9. Support Vector Machines (SVM)

Margin: $\frac{2}{\|w\|}$

Constraints:

$w^T x_i + b \geq 1$ for $y_i = +1$

$w^T x_i + b \leq -1$ for $y_i = -1$

Thus: $y_i(w^T x_i + b) - 1 \geq 0$ for $\forall x_i$

Optimization Problem:

Minimizer: $\frac{1}{2} w^T w$

Subject to: $f_i(w, b) = y_i(w^T x_i + b) - 1 \geq 0$

Lagrangian Dual Function:

Dual function: $g(\alpha) = \min_{\theta \in \mathbb{R}^d} (f_0(\theta) + \sum_{i=1}^M \alpha_i f_i(\theta))$

Lagrangian: $L(\theta, \alpha) = f_0(\theta) + \sum_{i=1}^M \alpha_i f_i(\theta)$

Conditions: $\alpha_i \geq 0$ and $f_i(\theta) \leq 0$

SVM Optimization Steps

1. Calculate Lagrangian:

$L(w, b, \alpha) = \frac{1}{2} w^T w - \sum_i \alpha_i [y_i(w^T x_i + b) - 1]$

2. Minimize L:

$\frac{\partial L}{\partial w} = w - \sum_i \alpha_i y_i x_i = 0$

$\frac{\partial L}{\partial b} = \sum_i \alpha_i y_i = 0$

$\Rightarrow w = \sum_i \alpha_i y_i x_i$

3. Dual Problem:

$g(\alpha) = \sum_i \alpha_i - \frac{1}{2} \sum_i \sum_j y_i y_j \alpha_i \alpha_j x_i^T x_j$

Note: $x_i^T x_j$ can be replaced by Kernel $\Phi(x_i, x_j)$

w.r.t. $\alpha_i \geq 0$, $\sum_i \alpha_i y_i = 0$

$w = \sum_i \alpha_i^* y_i x_i$

$b = \frac{1}{y_i} - w^T x_i = y_i - w^T x_i$

$\therefore h(x) = \text{sign}(\sum_{i \in S} \alpha_i y_i x_i^T x + b)$

Soft SVM (relaxed margin):

Constraint: $y_i(w^T x_i + b) \geq 1 - \xi_i$ ($\xi_i \geq 0$)

Minimize: $f_0(w, b, \xi) = \frac{1}{2} w^T w + C \sum_i \xi_i$

w.r.t.:

$y_i(w^T x_i + b) - 1 + \xi_i \geq 0$

$\xi_i \geq 0$

$(C \rightarrow \infty)$: hard margin

Soft Margin SVM Derivation

1. Lagrangian Formulation:

$L(w, b, \xi, \alpha) = \frac{1}{2} w^T w + C \sum_i \xi_i - \sum_i \alpha_i [y_i(w^T x_i + b) - 1 + \xi_i] - \sum_i \mu_i$

2. Minimize L (Gradients):

$\frac{\partial L}{\partial w} = w - \sum_i \alpha_i y_i x_i = 0$

$\frac{\partial L}{\partial b} = \sum_i \alpha_i y_i = 0$

$\frac{\partial L}{\partial \xi_i} = C - \alpha_i - \mu_i = 0 \Rightarrow \alpha_i = C - \mu_i$

Since $\mu_i \geq 0$ and $\alpha_i \geq 0$: Box Constraint $\alpha_i \in [0, C]$

3. Maximize Dual Function:

$g(\alpha) = \sum_i \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j x_i^T x_j$

Subject to:

$\sum_i \alpha_i y_i = 0$

$0 \leq \alpha_i \leq C$

Interpretation of α_i :

If $0 < \alpha_i < C$: $\xi_i = 0$ (point exactly on margin)

If $\alpha_i = C$: $\xi_i > 0$ (point violates margin)

Larger $C \rightarrow$ less tolerance for points inside margin

Hinge Loss:

$\frac{1}{2} w^T w + C \sum_i \max\{0, 1 - y_i(w^T x_i + b)\}$

Kernel Methods

Definition: $k(x_i, x_j) = \phi(x_i)^T \phi(x_j)$

Prediction: $h(x) = \text{sign}(\sum_{j \in S} \alpha_j y_j k(x_i, x_j) + b)$

Kernel Matrix Properties: Must be symmetric positive semi-definite

Valid Kernel Construction Rules:

1. Sum: $k(x_1, x_2) = k_1 + k_2$

2. Scaling: $k(x_1, x_2) = c k_1$ with $c > 0$

3. Product: $k(x_1, x_2) = k_1 \cdot k_2$

4. Transformation: $k(x_1, x_2) = k_3(\phi(x_1), \phi(x_2))$

5. Matrix Scaling: $k(x_1, x_2) = x_1^T A x_2$ where A is symmetric positive semi-definite

Common Examples:

Polynomial: $k(a, b) = (a^T b)^n$ or $(a^T b + 1)^p$

Gaussian (RBF): $k(a, b) = \exp\left(-\frac{\|a-b\|^2}{2\sigma^2}\right)$

Multiclass Classification Strategies

1 vs n classification: Look at maximum distance

1 vs 1 classification: Look at majority vote

10. Dimension Reduction PCA SVD

Dimension Reduction (PCA)

Transformation: $\tilde{x}' = \tilde{x} \cdot F \quad \Sigma_{x'} = F^T \Sigma_x F$

1. Centering Data

$$\tilde{x}_i = x_i - \bar{x} \text{ where } \bar{x} = \begin{bmatrix} \bar{x}_1 \\ \vdots \\ \bar{x}_d \end{bmatrix} = \frac{1}{N} \cdot X^T \cdot 1_N$$

2. Variance and Covariance

Variance: $\text{Var}(X_j) = \frac{1}{N} X_j^T X_j - \bar{x}_j^2$

Covariance: $\text{Cov}(X_i, X_j) = \frac{1}{N} X_i^T X_j - \bar{x}_i \bar{x}_j$

Covariance Matrix: $\Sigma_{\tilde{x}} = \frac{1}{N} \tilde{X}^T \tilde{X}$ (symmetric)

3. Eigen-decomposition

$\Sigma_{\tilde{x}} = \Gamma \Lambda \Gamma^T$ (where Λ is diagonal)

4. Transformation

$\cdot Y = \tilde{X} \cdot \Gamma$

$\cdot Y_{\text{reduced}} = \tilde{X} \cdot \Gamma_{\text{truncated}}$

Criteria: Retain 90% of variance: $\sum_{i=1}^k \lambda_i \geq 0.9 \sum_{i=1}^d \lambda_i$

Complexity: $O(nd^2 + d^3)$

5. Iterative Eigenvector Calculation

Power Iteration: $v \leftarrow \frac{v}{\|v\|}$ (converges to eigenvector with largest eigenvalue)

Singular Value Decomposition (SVD)

Goal: Find best low-rank approximation of matrix A

Frobenius Norm Objective: $\|A - B\|_F^2 = \sum_i \sum_j (a_{ij} - b_{ij})^2$

Complexity: $O(n \cdot d^2)$ or $O(n^2 \cdot d)$

Decomposition: $A = U \Sigma V^T$ where:

$\cdot U \in \mathbb{R}^{r \times r}$ (user-to-concept similarity)

$\cdot \Sigma \in \mathbb{R}^{r \times r}$ (item-to-concept similarity)

$\cdot V \in \mathbb{R}^{d \times r}$ (Rank-2 Decomposition: $A = (\sigma_1 u_1 v_1^T) + (\sigma_2 u_2 v_2^T) + \dots$)

Using SVD for Dimensionality Reduction

Projection: $P = U \Sigma$ or $P = A \cdot V^T$

Retain 90% energy: $\sum_{i=1}^r \sigma_i^2 \geq 0.9 \sum_{i=1}^d \sigma_i^2$

Relationship to Eigenvectors:

$\cdot V$ contains eigenvectors of $X^T X$

$\cdot U$ contains eigenvectors of $X X^T$

Matrix Factorization (MF)

1. Fundamentals & Metrics

RMSE: $\text{RMSE} = \sqrt{\frac{1}{|S|} \sum (r_{ui} - \hat{r}_{ui})^2}$

SSE: $\text{SSE} = \sum (r_{ui} - [\Sigma \Sigma V^T]_{ui})^2$

Decomposition: $R = U \Sigma V^T \approx Q \cdot P^T$

Prediction: $\hat{r}_{ui} = q_u \cdot p_i^T$

2. Alternating Optimization

1. Initialize: $P_0, Q_0, t = 0$

2. Update P : $P_{t+1} = \text{argmin}_P f(P, Q_t)$

Closed form: $p_i^T = \left(\frac{1}{|S_{*,i}|} \sum q_u^T q_u \right)^{-1} \cdot \frac{1}{|S_{*,i}|} \sum q_u^T r_{ui}$

3. Update Q : $Q_{t+1} = \text{argmin}_Q f(P_{t+1}, Q)$

4. Repeat until convergence

3. Stochastic Gradient Descent (SGD)

$\cdot e_{ui} \leftarrow r_{ui} - q_u \cdot p_i^T$

$\cdot q_u \leftarrow q_u + 2r(e_{ui} p_i)$

$\cdot p_i \leftarrow p_i + 2r(e_{ui} q_u)$ (r : learning rate)

4. Extensions: Bias & Regularization

Regularized Objective:

$\sum (r_{ui} - q_u \cdot p_i^T)^2 + \lambda_1 \sum \|q_u\|^2 + \lambda_2 \sum \|p_i\|^2$

Full Loss with Bias:

$L = \sum (r_{ui} - (q_u p_i^T + b_u + b_i + b))^2 + \lambda_1 \sum \|q_u\|^2 + \lambda_2 \sum \|p_i\|^2$

SGD Updates (with Bias & Regularization):

1. $e_{ui} = r_{ui} - q_u \cdot p_i^T$

2. $q_u = q_u + 2r(e_{ui} p_i - \lambda_2 p_i)$

3. $p_i = p_i + 2r(e_{ui} q_u - \lambda_2 p_i)$

4. $b_u = b_u + 2r e_{ui}$

5. $b_i = b_i + 2r e_{ui}$

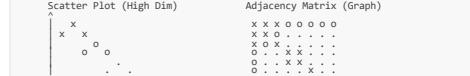
6. $b = \frac{1}{|S|} \sum r_{ui}$ (Global Bias)

11. Dimension Reduction Neighbor Graph Method

Neighbor Graph Method

Preserve local structure

Plaintext



t-SNE

High-dimensional similarity:

$$p_{j|i} = \frac{\exp(-\frac{\|x_i - x_j\|^2}{2\sigma_i^2})}{\sum_{k \neq i} \exp(-\frac{\|x_i - x_k\|^2}{2\sigma_i^2})}$$

Data X

$$p_{ii} = 0, p_{ij} = \frac{p_{ji} + p_{ii}}{2N}$$

$p_{ij} = p_{ji}$

Low-dimensional similarity:

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_k \sum_{l \neq k} (1 + \|y_k - y_l\|^2)^{-1}}$$

Target y

$$q_{ii} = 0$$

Next: minimize KL Divergence

$$\sum_i \sum_{j \neq i} p_{ij} \log \frac{p_{ij}}{q_{ij}} = KL(P||Q)$$

$\cdot KL(P||Q) \geq 0, = 0 \text{ iff } P = Q$

Divergence Behavior:

Plaintext



Plaintext



Autoencoder

Plaintext



Notes:

Extract key features and reconstruct; equivalent to only taking rank(X) = L.

Formula: $X W_1 W_2 = \tilde{X}$

$W_1 W_2 = W^* = \Gamma$

Γ is the largest L eigenvectors of $X^T X$.

12. Clustering

K-means: Objective (Distortion Measure)

Formula:

$$J(X, Z, \mu) = \sum_i^K \sum_k z_{ik} \|x_i - \mu_k\|^2$$

Indicator Variable:

$z_{ik} = \{0, 1, 0, 0, 0\}$ implies sample i belongs to class 3.

Lloyd's Algorithm (Alternating Optimization)

1. Initialize: All centroids μ_i .

2. Assignment: Assign cluster indicators based on the nearest neighbor.

... (Note: Whichever centroid a point is closest to, it belongs to that cluster.)

3. Update:

$$\mu_k = \frac{1}{N_k} \sum_i^K z_{ik} x_i$$

Where $N_k = \sum_i^K z_{ik}$

4. Loop: Repeat until convergence.

K-means ++

1. Randomly select one data point as μ_1 .

2. Calculate the distance of the remaining points to it: $\|x_i - \mu_1\|^2$.

3. Sample the next center point, with probability proportional to the distance size.

4. Re-calculate $D_i^2 = \min\{\|x_i - \mu_1\|^2, \|x_i - \mu_2\|^2, \dots\}$.

5. Repeat steps 3 and 4 until K centers are chosen.

Gaussian Mixture Model (GMM)

Model Definition:

$$p(x|\theta) = \sum_i p(x|z_i, \theta) p(z_i|\theta)$$

z: latent variables

Optimization Goal: $\theta^* = \text{argmax}_{\theta} p(x|\theta)$

Distributions:

1. Prior: $p(z|\theta) = \text{Cat}(\pi) // \{\pi_1, \pi_2, \dots, \pi_K\}$

2. Constraint: $\sum_i^K \pi_i = 1$

2. Likelihood: $p(x|z = k, \theta) = \mathcal{N}(x|\mu_k, \Sigma_k)$

Parameters:

Thus $\theta = \{\pi, \mu, \Sigma\}$

• π : k parameters (effectively k - 1)

• $\mu \in \mathbb{R}^d$: k · d parameters

• $\Sigma \in \mathbb{R}^{d \times d}$: k · $\frac{d(d+1)}{2}$ due to symmetry)

Using GMM as a Generative Model

1. Sample class according to $\pi = \{\pi_1, \pi_2, \dots, \pi_K\}$

2. Generate x according to $x \sim \mathcal{N}(\mu_k, \Sigma_k)$ (Probability Density)

Marginal Probability:

$$p(x|\pi, \mu, \Sigma) = \sum_i^K \pi_i \cdot \mathcal{N}(x|\mu_k, \Sigma_k)$$

Log-Likelihood:

$$\log p(X|\pi, \mu, \Sigma) = \sum_i^K \log \left(\sum_k^K \pi_k \cdot \mathcal{N}(x_i|\mu_k, \Sigma_k) \right)$$

Inference

Calculating the responsibility of component k for observation i:

$$\pi_{ik} = \frac{\pi_k \mathcal{N}(x_i|\mu_k, \Sigma_k)}{\sum_j^K \pi_j \mathcal{N}(x_i|\mu_j, \Sigma_j)}$$

Notation: